

Bimodal Covariates and the Estimated MCCE

Brenton Kenkel — PSCI 8357

February 11, 2016

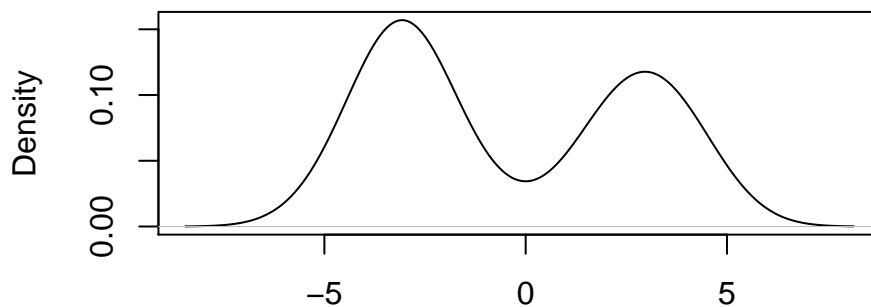
In class today, I mentioned that the estimated standard error of the MCCE for a quadratic relationship tends to be smallest for values close to the sample mean. This raised the question of whether this holds true with a bimodal covariate—if there’s relatively little data near the sample mean, is the standard error of the estimated MCCE lowest there? This note will run a little simulation to find out.

First we need to know how to simulate a bimodal covariate. Assume there is a “left” distribution with mean -3 and a “right” distribution with mean 3 , and each observation has a 50-50 chance of being chosen from the left or the right.

```
n_obs <- 100
x_left <- rnorm(n_obs, mean = -3)
x_right <- rnorm(n_obs, mean = 3)
is_left <- rbinom(n_obs, size = 1, prob = 0.5)
x <- is_left * x_left + (1 - is_left) * x_right
```

Let’s check that the resulting distribution is indeed bimodal.

```
plot(density(x), xlab = "", main = "")
```

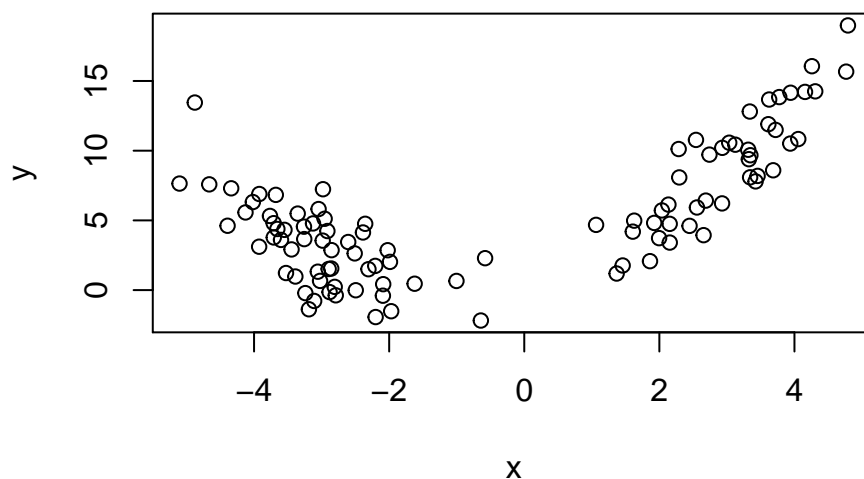


OK, so now let’s generate a response variable according to a quadratic model:

$$Y = 1 + \beta_1 X + \beta_2 X^2 + \epsilon.$$

As in class, we'll assume that the expected response is greater at the extremes (e.g., people with extreme viewpoints are more likely to vote). We'll also assume a slight positive trend overall (e.g., conservatives are more likely to vote, all else equal).

```
y <- 1 + x + 0.5 * x^2 + rnorm(n_obs, sd = 2)
plot(x, y)
```



We can fit a quadratic model and use it to estimate the pointwise marginal changes in conditional expectation,

$$MCCE(X) = \beta_1 + 2\beta_2 X.$$

```
fit_quadratic <- lm(y ~ x + I(x^2))
summary(fit_quadratic)
```

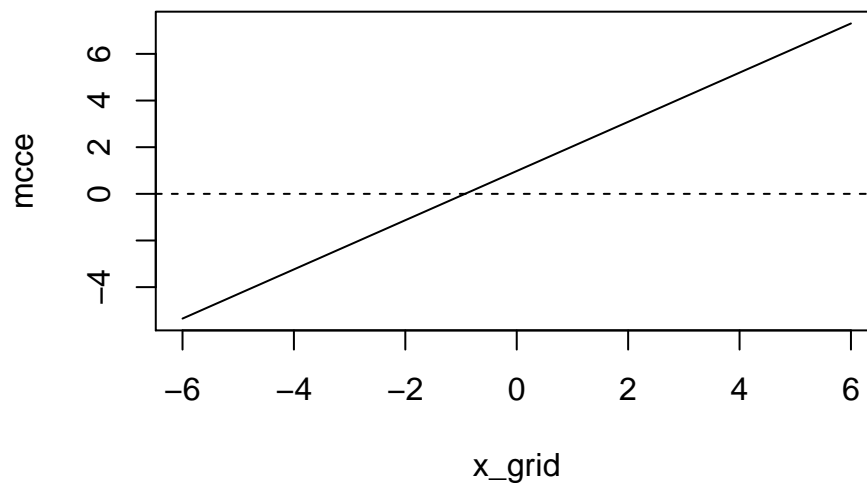
```
##
## Call:
## lm(formula = y ~ x + I(x^2))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.396 -1.858  0.153  1.477  4.905
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   0.7792     0.4163   1.87   0.064
```

```
## x          0.9781    0.0659   14.85   <2e-16
## I(x^2)     0.5272    0.0371   14.20   <2e-16
##
## Residual standard error: 2.04 on 97 degrees of freedom
## Multiple R-squared:  0.804, Adjusted R-squared:  0.8
## F-statistic: 199 on 2 and 97 DF,  p-value: <2e-16
```

To estimate the MCCE function, we'll calculate the estimated MCCE for each point in a grid of potential values of X .

```
x_grid <- seq(-6, 6, length.out = 100)
beta_1 <- coef(fit_quadratic)["x"]
beta_2 <- coef(fit_quadratic)["I(x^2)"]
mcce <- beta_1 + 2 * beta_2 * x_grid

plot(x_grid, mcce, type = "l")
abline(h = 0, lty = 2)
```



At this point we could use `vcov(fit_quadratic)` to estimate the standard error of the MCCEs. But we're not interested in the estimated standard errors—we want the *true* standard errors. In other words, is the estimated MCCE more variable at the extremes (where more of the data lies) or near the mean of X (where data is scarce)? To do this, we'll simulate the whole process repeatedly.

```
library("foreach")
```

```

n_obs <- 100
x_grid <- seq(-6, 6, length.out = 100)

sim_results <- foreach (i = 1:1000, .combine = "rbind") %do% {
  ## Simulate data
  x_left <- rnorm(n_obs, mean = -3)
  x_right <- rnorm(n_obs, mean = 3)
  is_left <- rbinom(n_obs, size = 1, prob = 0.5)
  x <- is_left * x_left + (1 - is_left) * x_right
  y <- 1 + x + 0.5 * x^2 + rnorm(n_obs, sd = 2)

  ## Fit regression
  fit_quadratic <- lm(y ~ x + I(x^2))

  ## Extract coefficients and calculate MCCE for each element
  ## in the grid
  beta_1 <- coef(fit_quadratic)["x"]
  beta_2 <- coef(fit_quadratic)["I(x^2)"]
  mcce <- beta_1 + 2 * beta_2 * x_grid

  ## Return pointwise MCCEs
  mcce
}

```

Let's take a look at the simulation results.

```
dim(sim_results)
```

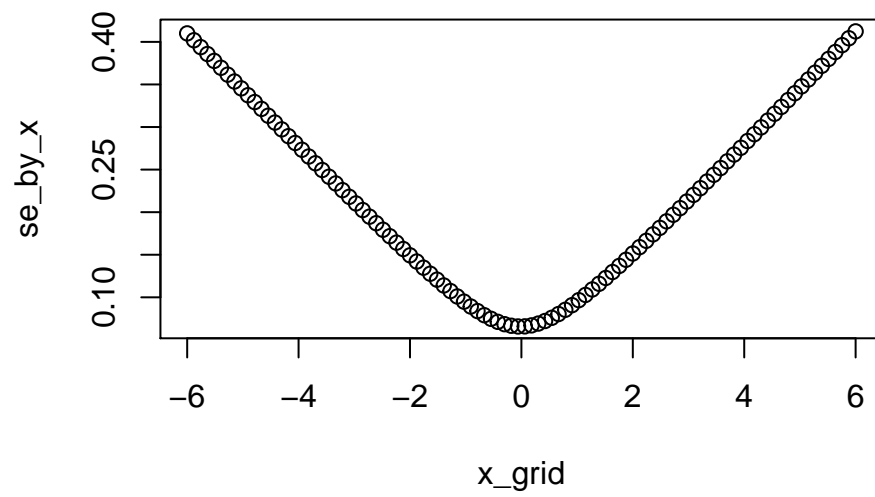
```
## [1] 1000 100
```

Each column is a value of `x_grid`, our grid of potential values of X . Each row is the estimated set of MCCEs from one iteration of the simulation. So to get the standard error corresponding to each value of X , we will take the standard deviation of each column of `sim_results`. We use `apply()` to apply the same function to each column of a matrix.

```

se_by_x <- apply(sim_results,
  MARGIN = 2, # 1 for rows, 2 for columns
  FUN = sd)
plot(x_grid, se_by_x)

```



So we see that the standard errors—the true ones, not just the estimated ones—are lowest near the mean of X , even though there is relatively little data there.